

# TP ON ADVERSARIAL BANDITS

In the lecture, we saw the adversarial bandit framework as a game between a player and nature. In fact, there is a strong connexion between regret minimization and game theory. In this practical session, we will apply the EXP3 algorithm to a sequential two-player zero sum game.

We consider the sequential version of a two-player zero-sum games between a player and an adversary.

Let  $L \in [-1, 1]^{M \times N}$  be a loss matrix.

At each round  $t = 1, \dots, T$

- The player choose a distribution  $p_t \in \Delta_M := \{p \in [0, 1]^M, \sum_{i=1}^M p_i = 1\}$
- The adversary chooses a distribution  $q_t \in \Delta_N$
- The actions of both players are sampled  $i_t \sim p_t$  and  $j_t \sim q_t$
- The player incurs the loss  $L(i_t, j_t)$  and the adversary the loss  $-L(i_t, j_t)$ .

Setting 1: Setting of a sequential two-player zero sum game

1. Define  $M, N$  and a loss matrix  $L \in [-1, 1]^{M \times N}$  that corresponds to the game “Rock paper scissors”<sup>1</sup>.

## Full information feedback

In this part, we assume that both players know the matrix  $L$  in advance and can compute  $L(i, j)$  for any  $(i, j)$ .

2. Implementation of EWA.

- (a) In order to implement the exponential weight algorithm, you need a way to sample from the exponential weight distribution. Implement the function `rand_exp` that takes as input a probability vector  $p \in \Delta_M$  and uses a single call to `rand()` to return  $X \in [M]$  with  $P(X = i) = p_i$ .
- (b) Define a function `EWA_update` that takes as input a vector  $p_t \in \Delta_M$  and a loss vector  $\ell_t \in [-1, 1]^M$  and return the updated vector  $p_{t+1} \in \Delta_M$  defined for all  $i \in [M]$  by

$$p_{t+1}(i) = \frac{p_t(i) \exp(-\eta \ell_t(i))}{\sum_{j=1}^M p_t(j) \exp(-\eta \ell_t(j))}.$$

3. Simulation against a fixed adversary. Consider the game “Rock paper scissors” and assume that the adversary chooses  $q_t = (1/2, 1/4, 1/4)$  and samples  $j_t \sim q_t$  for all rounds  $t \geq 1$ .
  - (a) What is the loss  $\ell_t(i)$  incurred by the player if he chooses action  $i$  at time  $t$ ? Simulate an instance of the game for  $t = 1, \dots, T = 100$  for  $\eta = 1$ .

<sup>1</sup>This is a common game where two players choose one of 3 options: (Rock, Paper, Scissors). The winner is decided according to the following: Rock crushes scissors, Paper covers Rock, Scissors cuts paper

- (b) Plot the evolution of the weight vectors  $p_1, p_2, \dots, p_T$ . What seems to be the best strategy against this adversary?
  - (c) Plot the average loss  $\bar{\ell}_t = \frac{1}{t} \sum_{s=1}^t \ell(i_s, j_s)$  as a function of  $t$ .
  - (d) Plot the cumulative regret.
  - (e) To see if the algorithm is stable, repeat the simulation  $n = 10$  times and plot the average loss  $(\bar{\ell}_t)_{t \geq 1}$  obtained in average, in maximum and in minimum over the  $n$  simulations.
  - (f) Repeat one simulation for different values of learning rates  $\eta \in \{0.01, 0.05, 0.1, 0.5, 1\}$  and plot the final regret as a function of  $\eta$ . What are the best  $\eta$  in practice and in theory.
4. Simulation against an adaptive adversary. Repeat the simulation of question 3) when the adversary is also playing EWA with learning parameters  $\eta = 0.05$ .
- (a) Plot  $\frac{1}{t} \sum_{s=1}^t \ell(i_s, j_s)$  as a function of  $t$ .

It is possible to show that if both players play according to a regret minimizing strategy the cumulative loss of the player converges to the value of the game

$$V = \min_{p \in \Delta_M} \max_{q \in \Delta_q} p^\top Lq.$$

- (b) Define  $\bar{p}_t = \frac{1}{t} \sum_{s=1}^t p_s$ . Plot in log log scale  $\|\bar{p}_t - (1/3, 1/3, 1/3)\|_2$  as a function of  $t$ .

It is possible to show that  $(\bar{p}_t, \bar{q}_t)_{t \geq 1}$  converges almost surely to a Nash equilibrium of the game. This means that if  $p \times q$  is a Nash equilibrium, none of the players should change its strategy if the other player does not change hers.

## Bandit feedback

Now, we assume that the players do not know the game in advance but only observe the performance  $L(i_t, j_t)$  (that we assume here to be in  $[0, 1]$ ) of the actions played at time  $t$ . They need to learn the game and adapt to the adversary as one goes along.

- 5. Implementation of EXP3. Since both players are symmetric, we focus on the first player.
  - (a) Implement the function `estimated_loss` that takes as input the action  $i_t \in [M]$  played at round  $t \geq 1$  and the loss  $L(i_t, j_t)$  suffered by the player and return the vector of estimated loss  $\hat{\ell}_t \in \mathbb{R}_+^M$  used by EXP3.
  - (b) Implement the function `EXP3_update` that takes as input a vector  $p_t \in \Delta_M$ , the action  $i_t \in [M]$  played by the player and the loss  $L(i_t, j_t)$  and return the updated weight vector  $p_{t+1} \in \Delta_M$ .
- 6. Repeat Questions 3.a) to 3.f) with EXP3 instead of EWA.
- 7. Repeat Question 4.a) and 4.b) with EXP3 instead of EWA.

## Optional extensions

- 8. Repeat Question 4.a) when the adversary is playing a UCB algorithm. Who wins between UCB and EXP3?

9. In this lecture, we saw that EXP3 has a sublinear expected regret. Yet, as shown by question 6.e), it is extremely unstable with a large variance. Implement **EXP3.IX** (see Chapter 12 of [1]) a modification of **EXP3** that controls the regret in expectation and simultaneously keeps it stable. Repeat question 3.e) with **EXP3.IX**
10. Try different games (not necessarily zero-sum games). In particular, how these algorithms behave for the prisoner's dilemma (see wikipedia)? The prisoner's dilemma is a two-player games that shows why two completely rational individuals might not cooperate, even if it appears that it is in their best interests to do so. The losses matrices are:

$$L^{(player)} = \begin{pmatrix} 1 & 3 \\ 0 & 2 \end{pmatrix} \quad \text{and} \quad L^{(adversary)} = \begin{pmatrix} 1 & 0 \\ 3 & 2 \end{pmatrix}.$$

## References

- [1] Tor Lattimore and Csaba Szepesvári. Bandit algorithms. <https://tor-lattimore.com/downloads/book/book.pdf>, 2019.